

# How this was put together

A method note for the general reader

Companion to *The price of a probability*  
Ethan Gallant · September 2026

---

This note explains where the numbers came from, how they were checked, and how the judgements were reached. **No prior knowledge of banking is assumed.** It also records what went wrong, because a method note that only lists successes is not a method note.

## 1. The question

---

**When a war breaks out, what does it actually cost, and who can be shown to have paid it?**

Most writing about geopolitical risk stops at description: this is serious, this could escalate, watch this space. It rarely lands on a number that somebody has had to put in writing and stand behind.

**Banks are the exception, and that is why this report is about banks.**

## 2. Why a bank is the place to look

---

A bank lends money. Some of it will not come back, and accounting rules require the bank to set aside an estimate of those future losses *now*, before anyone has missed a payment.

To make that estimate it has to describe the future. It does so with a small set of scenarios — roughly a good world, a normal world and a bad world — and assigns each a probability. **Those probabilities must add up to 100 per cent.**

---

**So if a bank decides the world has become more dangerous, it has to take probability away from the normal world and give it to the bad one. That reassignment costs real money, immediately, in the accounts.**

---

Two more terms, and then the jargon is finished:

- **A management overlay** is money set aside *on top* of what the computer model produced, because management thinks the model is missing something. It is a judgement, written down as a number.
- **Staging** is the bank's grading of how healthy a loan is — stage 1 performing, stage 2 deteriorating, stage 3 in trouble. Moving a loan down the stages costs money too.

**The point of all this: a view about a war becomes a figure in a published set of accounts.** That is an unusually hard piece of evidence, and it is public.

## 3. How the numbers were actually found

---

### Standard Chartered

The starting point was the bank's own results press release. It showed a total credit impairment charge of \$446 million, with a note that part of it was "mostly driven by management overlays relating to the Middle East conflict".

**That was suggestive but not usable:** a press release does not break the figure down. The breakdown sits in a separate regulatory announcement — the half-year results Part 2 filing, roughly a hundred thousand characters that almost nobody reads — where a passage attributes **\$234 million** to the conflict and splits it three ways.

**The split adds up exactly: \$159m + \$14m + \$61m = \$234m.** That mattered: it meant the bank had given a complete account rather than a headline, so the figure could be used without reconstruction.

### First Abu Dhabi Bank, and a dead end

For comparison, a bank inside the region. First Abu Dhabi Bank's own website was down — every document returned an empty file.

**The filings were obtained instead from the Abu Dhabi Securities Exchange**, where the bank is legally required to lodge them — not a workaround or a mirror, but the primary source, and arguably a better one, because it is the version filed with the regulator.

### Two more banks, to test whether two was a pattern

Standard Chartered and First Abu Dhabi Bank had moved by almost identical amounts, at almost identical cost relative to profit. **That looked like a finding.** Emirates NBD and HSBC were added specifically to test it, and it did not survive: Emirates NBD went further than either in the first quarter and then unwound the whole thing by June, and HSBC declined to put a figure on it at all.

---

**The symmetry was a coincidence of two data points. Testing it destroyed it, and the real finding is what replaced it: four banks looked at the same war and reached four different answers.**

---

### The war itself

Conflict facts come from three **House of Commons Library research briefings**, written by parliamentary researchers for MPs of all parties and non-partisan by design. Using them instead of news coverage removes an argument about sources before it starts.

### The chokepoint

The claim the report is built on — that **most of the oil going through the Strait of Hormuz has no other way out of the region** — comes from the US Energy Information Administration, a government statistical agency, in those words.

**An earlier draft turned that into a number: 87 per cent.** Twenty million barrels a day of normal flow, minus 2.6 million of pipeline capacity, equals 87 per cent stranded. It reads well and it is wrong: **the 2.6 million is spare capacity available in a crisis, not the total volume that has a route.** The pipelines are bigger than that, but already partly in use.

**The fix was to quote the agency and state the capacity as capacity** — weaker-looking and stronger, because the claim now belongs to the EIA rather than to an arithmetic step of Ethan's own.

### What could not be obtained

Credit rating agency research sits behind paywalls or on pages that would not load. **Rather than cite a rating second-hand, the report was restructured so it does not depend on any rating claim at all**, and says so on its final page.

## 4. The rule that governed everything

---

---

**A search engine summary is a lead, never a source.**

---

This rule exists because it was broken. An early version of the working file described every figure as traced to a primary source. **That was not true** — several had come from search result summaries of documents nobody had opened.

The fix was to go through the file line by line and tag every figure either **PRIMARY** — read in the bank’s own filing — or **UNVERIFIED**, with a standing rule that nothing in the second category could be published.

**It caught something.** A figure for Emirates NBD’s first-quarter charge had been carried for days on secondary reporting. It survived into a late draft of the report, including the summary table on the front page, and was only removed at the final check. **The argument turned out not to need it.** What replaced it is a sentence saying the bank does not disclose that figure at all — which is more honest and, as it happens, more interesting.

## 5. How the judgement was made

---

The report’s central claim is that Standard Chartered’s 30 per cent weighting on a sustained Middle East conflict is now too low, and that the right figure is **80 to 90 per cent**. **One caveat is stated up front:** a bank’s scenario weighting and an intelligence probability are not the same kind of object, so the comparison is a translation, not an equivalence.

### Why those words and those numbers

Probabilities use the **PHIA yardstick**, the scale used across UK government intelligence assessment. It fixes what words mean: “unlikely” is 25 to 35 per cent, “highly likely” is 80 to 90. **The scale has deliberate gaps between the bands** — you cannot say 45 to 55 per cent, because that is where an assessment stops saying anything. Every judgement here sits inside a single band.

### Probability and confidence are different things

**Probability** is how likely the thing is. **Confidence** is how good the evidence behind that estimate is. They are stated separately throughout, because a well-evidenced 60 per cent and a guessed 60 per cent are not the same claim.

The central judgement is **80 to 90 per cent, confidence medium**: high confidence on the physical mechanics, because the strait is shut and there is no way round it; lower confidence on how long the politics will last.

### What the judgement rests on

Three things happened in August, after all four banks had closed their books in June:

1. Iran appointed a hardliner to run national security, and his stated conditions for reopening the strait are *wider* than what was on the table in June.
2. **The UAE announced it would cut off all trade with Iran** — its largest single source of imports.
3. The United States announced a greatly expanded sanctions regime.

**The reasoning in one line: a state under severe economic pressure that raises its asking price is not preparing to settle.**

### The distinction that keeps this honest

**The claim is not that the banks were wrong in June.** On what was knowable at the time, their judgement was defensible. The claim is that it is not defensible now.

That is the difference between hindsight and assessment, and it is also the more useful thing to say, because it is about what happens next rather than who was wrong.

## 6. Setting the things to watch

---

The report ends with indicators — specific, numbered conditions that would show the assessment coming true or falling apart.

**These were written before the current readings were looked up** — the only way to stop a threshold being quietly set wherever the data already happens to be.

**It went wrong once.** An early draft set a war-risk insurance trigger at 3 per cent of a ship’s value; the market had been between 3 and 8 per cent for months. The threshold described the present rather than predicting anything, and was reset. **A threshold that is already breached is not a threshold**, and the error is recorded rather than quietly corrected.

Two further tests were applied to each indicator: **can I actually observe this from public sources**, and **does it lead or merely confirm**? One indicator was cut for failing the first. Rating actions are kept but labelled as confirming, because agencies move after the damage is visible.

## 7. What was cut, and why

| What                                                            | Why it went                                                                                                     |
|-----------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------|
| Near-identical cost as a share of profit at two banks           | A coincidence of two data points; two more banks did not fit                                                    |
| A claim that credit quality improved                            | The evidence showed only that backward-looking indicators had not worsened at the same time. Not the same thing |
| An Emirates NBD figure carried in secondary reporting           | Never traced to the bank's own filing                                                                           |
| Three Standard Chartered figures, and a comparison against 2024 | Not found in the documents actually read                                                                        |
| Any claim about credit ratings                                  | Could not be verified at source                                                                                 |
| A fabricated citation                                           | A web search link dressed up as a research institute's work; it cited nothing                                   |
| Several confident verbs                                         | Claimed more than the filings showed                                                                            |

**The project also changed shape once.** It began drifting into a study of bank disclosure practice — technically correct, of no interest outside an accounting department — and was rebuilt to start with the war and work forwards to who pays.

**A last pass asked a different system to attack the finished report as a hostile reader.** Ten corrections came back; nine were accepted, including the 87 per cent above. **One was rejected** — a challenge to a trade statistic, traced to WTO data cited by the Commons Library and kept, with the attribution added so the question does not arise again. **Being told you are wrong is cheap. Checking which half is right is the work.**

## 8. The moment the research changed the argument

The report was expected to argue the banks were modelling the wrong thing. Earlier work on the 2024 Red Sea disruption had found the damage ran through rerouting, insurance and freight cost rather than lost volume: cargo still arrived, later and dearer.

**The chokepoint numbers broke that expectation.** The Red Sea had an alternative route round the Cape of Good Hope. The Strait of Hormuz has enough pipeline capacity for roughly an eighth of what normally passes through it.

**So the banks may well be modelling the right thing after all, and the useful finding is the one underneath: chokepoints are not interchangeable.** A scenario built by analogy to the last crisis would have modelled the wrong channel entirely.

**That section survives because the prior was wrong, not despite it.**

## 9. What would show this wrong, and what is still open

The report lists five conditions that would overturn it: Hormuz reopening, Standard Chartered releasing the money it set aside, Emirates NBD not needing any more, a negotiated settlement, and the bank's next disclosure showing it modelled the other channel. **Each would overturn the judgement rather than soften it. An assessment that cannot be shown to be wrong is not an assessment.**

Two documents also remain unopened: Emirates NBD's first-quarter filing and First Abu Dhabi Bank's 2025 comparable. **Neither is load-bearing** — the first supports a figure that has been removed, the second a calculation that now carries its

caveat in the text. **Recorded rather than left out, because the standard the report sets for four banks applies to the report as well.**